

# **CS 5600 Computer Systems**

## **Spring 2026**

### **Virtualization: The CPU and Scheduling**

#### **Lecture 2**

#### **January 20, 2026**

Prof. Scott Valcourt  
[s.valcourt@northeastern.edu](mailto:s.valcourt@northeastern.edu)  
603-380-2860 (cell)

Portland, ME



# Lecture Agenda

- Abstraction – Chapter 4 – The Process
- Interlude – Chapter 5 – Process API
- Mechanism – Chapter 6 – Limited Direct Execution
- Scheduling – Chapter 7 – Introduction
- Scheduling – Chapter 8 – Multi-Level Feedback Queue
- Scheduling – Chapter 9 – Proportional Share
- Last Week: Chapter 1 - Introduction

# What is a Process?

Process: An **execution stream** in the context of a **process state**

What is an execution stream?

- Stream of executing instructions
- Running piece of code
- “thread of control”

What is process state?

- Everything that the running code can affect or be affected by
- Registers
  - General purpose, floating point, status, program counter, stack pointer
- Address space
  - Heap, stack, and code
- Open files

# **Processes vs. Programs**

A process is different than a program

- Program: Static code and static data
- Process: Dynamic instance of code and data

Can have multiple process instances of the same program

- Can have multiple processes of the same program  
Example: many users can run “ls” at the same time

# Process API

- These APIs are available on any modern OS.
  - **Create**
    - Create a new process to run a program
  - **Destroy**
    - Halt a runaway process
  - **Wait**
    - Wait for a process to stop running
  - **Miscellaneous Control**
    - Some kind of method to suspend a process and then resume it
  - **Status**
    - Get some status info about a process


# Process Creation

1. **Load** a program code into memory, into the address space of the process.
  - Programs initially reside on disk in *executable format*.
  - OS performs the loading process **lazily**.
    - Loading pieces of code or data only as they are needed during program execution.
2. The program's run-time **stack** is allocated.
  - Use the stack for *local variables*, *function parameters*, and *return address*.
  - Initialize the stack with arguments → argc and the argv array of main() function

## Process Creation (Cont.)

3. The program's **heap** is created.
  - Used for explicitly requested dynamically allocated data.
  - Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS does some other initialization tasks.
  - input/output (I/O) setup
  - Each process by default has three open file descriptors.
  - Standard input, output, and error
5. **Start the program** running at the entry point, namely `main()`.
  - The OS *transfers control* of the CPU to the newly-created process.


# Loading: From Program To Process



# Process States

- A process can be one of three states.
  - **Running**
    - A process is running on a processor.
  - **Ready**
    - A process is ready to run but for some reason the OS has chosen not to run it at this given moment.
  - **Blocked**
    - A process has performed some kind of operation.
    - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

# Process State Transition



# Data structures

- The OS has **some key data structures** that track various relevant pieces of information.
  - **Process list**
    - Ready processes
    - Blocked processes
    - Current running process
  - **Register context**
- PCB (Process Control Block)
  - A C-structure that contains information about each process.

# Example) The xv6 kernel Proc Structure

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;    // Index pointer register
    int esp;    // Stack pointer register
    int ebx;    // Called the base register
    int ecx;    // Called the counter register
    int edx;    // Called the data register
    int esi;    // Source index register
    int edi;    // Destination index register
    int ebp;    // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

# Example) The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;          // Start of process memory
    uint sz;            // Size of process memory
    char *kstack;       // Bottom of kernel stack
                        // for this process
    enum proc_state state; // Process state
    int pid;            // Process ID
    struct proc *parent; // Parent process
    void *chan;          // If non-zero, sleeping on chan
    int killed;          // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;   // Current directory
    struct context context; // Switch here to run process
    struct trapframe *tf; // Trap frame for the
                        // current interrupt
};
```

# How to provide the illusion of many CPUs?

- CPU virtualizing
  - The OS can promote the illusion that many virtual CPUs exist.
  - **Time sharing:** Running one process, then stopping it and running another
  - The potential cost is **performance**.

# Processes vs. Threads

- A process is different than a thread
- Thread: “Lightweight process” (LWP)
  - An execution stream that shares an address space
  - Multiple threads within a single process
- Example:
  - Two **processes** examining same memory address 0xffe84264 see **different** values (i.e., different contents)
  - Two **threads** examining memory address 0xffe84264 see **same** value (i.e., same contents)

# Virtualizing the CPU

## Goal:

Give each process the impression it alone is actively using CPU

Resources can be shared in **time** and **space**

Assume single uniprocessor

Time-sharing (multi-processors: advanced issue)

Memory?

Space-sharing (later)

Disk?

Space-sharing (later)

# How to Provide Good CPU Performance?

## Direct execution

- Allow user process to run directly on hardware
- OS creates process and transfers control to starting point (i.e., `main()`)

## Problems with direct execution?

1. Process could do something restricted
  - Could read/write other process data (disk or memory)
2. Process could run forever (slow, buggy, or malicious)
  - OS needs to be able to switch between processes
3. Process could do something slow (like I/O)
  - OS wants to use resources efficiently and switch CPU to other process

## Solution:

**Limited direct execution** – OS and hardware maintain some control

## **Problem 1: Restricted OPS**

How can we ensure user processes can't harm others?

Solution: privilege levels supported by hardware (bit of status)

- User processes run in user mode (restricted mode)
- OS runs in kernel mode (not restricted)
  - Instructions for interacting with devices
  - Could have many privilege levels (advanced topic)

How can process access device?

- System calls (function call implemented by OS)
- Change privilege level through system call (trap)

# System Call



P wants to call `read()`

# System Call



P can only see its own memory because of **user mode** (other areas, including kernel, are hidden)

# System Call



P wants to call `read()` but no way to call it directly


# System Call



read():

```
movl $6, %eax;    int $64
```

# System Call



read():

movl \$6, %eax;



syscall-table index


int \$64



trap-table index

# System Call

Kernel mode: we can do anything!



read():


```
movl $6, %eax; int $64
```

syscall-table index

trap-table index

# System Call

Follow entries to correct system call code



read():


movl \$6, %eax; int \$64

syscall-table index

trap-table index

# System Call

Kernel can access user memory to fill in user buffer  
return-from-trap at end to return to Process P



**read():**

```
movl $6, %eax;
```

1

## syscall-table index

int \$64

## trap-table index

# What to limit?

User processes are not allowed to perform:

- General memory access
- Disk I/O
- Special x86 instructions like `lidt` (load interrupt descriptor table)

What if process tries to do something restricted?

# Problem 2: How to take CPU AWAY?

OS requirements for **multiprogramming** (or multitasking)

- Mechanism
  - To switch between processes
- Policy
  - To decide which process to schedule when

Separation of policy and mechanism

- Reoccurring theme in OS
- **Policy:** Decision-maker to optimize some workload performance metric
  - Which process when?
  - Process **Scheduler:** Future lecture
- **Mechanism:** Low-level code that implements the decision
  - How?
  - Process **Dispatcher:** Today's lecture

# Dispatch Mechanism

OS runs **dispatch loop**

```
while (1) {  
    run process A for some time-slice  
    stop process A and save its context  
    load context of another process B  
}
```

Context-switch

Question 1: How does dispatcher gain control?


Question 2: What execution context must be saved and restored?

# **Q1: How does Dispatcher get CONTROL?**


## Option 1: Cooperative Multi-tasking

- Trust process to relinquish CPU to OS through traps
  - Examples:
    - System call
    - Page fault (access page not in main memory)
    - Error (illegal instruction or divide by zero)
  - Provide special `yield()` system call


# Cooperative Approach




# Cooperative Approach




# Cooperative Approach



# Cooperative Approach



# Cooperative Approach



# **Q1: How Does the Dispatcher RUN?**

- Problem with cooperative approach?
- Disadvantages: Processes can misbehave
  - By avoiding all traps and performing no I/O, can take over entire machine
  - Only solution: Reboot!
- Not performed in modern operating systems

# **Q1: How Does the Dispatcher RUN?**

## Option 2: True Multi-tasking

- Guarantee OS can obtain control periodically
- Enter OS by enabling periodic alarm clock
  - Hardware generates timer interrupt (CPU or separate chip)
  - Example: Every 10ms
- User must not be able to mask timer interrupt
- Dispatcher counts interrupts between context switches
  - Example: Waiting 20 timer ticks gives 200 ms time slice
  - Common time slices range from 10 ms to 200 ms

## **Q2: What Context must be Saved?**

Dispatcher must track context of process when not running

- Save context in **process control block (PCB)** (or, process descriptor)

What information is stored in PCB?

- PID
- Process state (I.e., running, ready, or blocked)
- Execution state (all registers, PC, stack ptr)
- Scheduling priority
- Accounting information (parent and child processes)
- Credentials (which resources can be accessed, owner)
- Pointers to other allocated resources (e.g., open files)

Requires special hardware support

- Hardware saves process PC and PSR on interrupts

# Context Switching

Operating System

Hardware

Program

Process A

...

# Context Switching

Operating System

Hardware

Program

Process A ...

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

# Context Switching

Operating System

Hardware

Program

Process A ...

Handle the trap  
Call `switch()` routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

# Context Switching

Operating System

Hardware

Program

Process A ...

Handle the trap  
Call `switch()` routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

# Context Switching

Operating System

Hardware

Program

Process A ...

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

Handle the trap  
Call `switch()` routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

Process B ...

# The xv6 Context Switch Code


```
1 # void swtch(struct context **old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl swtch
6 swtch:
7     # Save old registers
8     movl 4(%esp), %eax          # put old ptr into eax
9     popl 0(%eax)              # save the old IP
10    movl %esp, 4(%eax)         # and stack
11    movl %ebx, 8(%eax)         # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
18     # Load new registers
19    movl 4(%esp), %eax          # put new ptr into eax
20    movl 28(%eax), %ebp # restore other registers
21    movl 24(%eax), %edi
22    movl 20(%eax), %esi
23    movl 16(%eax), %edx
24    movl 12(%eax), %ecx
25    movl 8(%eax), %ebx
26    movl 4(%eax), %esp          # stack is switched here
27    pushl 0(%eax)              # return addr put in place
28    ret                      # finally return into new ctxt
```

## Problem 3: Slow Operations such as I/O?

When running process performs op that does not use CPU, OS switches to process that needs CPU (policy issues)

OS must track mode of each process:

- Running:
  - On the CPU (only one on a uniprocessor)
- Ready:
  - Waiting for the CPU
- Blocked
  - Asleep: Waiting for I/O or synchronization to complete



Transitions?

## **Problem 3: Slow Operations such as I/O?**

OS must track every process in system

Each process identified by unique Process ID (PID)

OS maintains queues of all processes

Ready queue: Contains all ready processes

Event queue: One logical queue per event

e.g., disk I/O and locks

Contains all processes waiting for that event to complete

Next Topic: Policy for determining which **ready** process to run

# Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
- OS handles these situations:
  - **Disable interrupts** during interrupt processing
  - Use several sophisticated **locking** schemes to protect concurrent access to internal data structures.

# Summary

Virtualization:

Context switching gives each process impression it has its own CPU

Direct execution makes processes fast

Limited execution at key points to ensure OS retains control

Hardware provides a lot of OS support

- user vs kernel mode
- timer interrupts
- automatic register saving

# **Break**

# Process Creation

Two ways to create a process

- Build a new empty process from scratch
- Copy an existing process and change it appropriately

Option 1: New process from scratch

- Steps
  - Load specified code and data into memory;  
Create empty call stack
  - Create and initialize PCB (make look like context-switch)
  - Put process on ready list
- Advantages: No wasted work
- Disadvantages: Difficult to setup process correctly and to express all possible options
  - Process permissions, where to write I/O, environment variables
  - Example: WindowsNT has call with 10 arguments

# Process Creation

## Option 2: Clone existing process and change

- Example: Unix fork() and exec()
  - Fork(): Clones calling process
  - Exec(char \*file): Overlays file image on calling process
- Fork()
  - Stop current process and save its state
  - Make copy of code, data, stack, and PCB
  - Add new PCB to ready list
  - Any changes needed to child process?
- Exec(char \*file)
  - Replace current data and code segments with those in specified file
- Advantages: Flexible, clean, simple
- Disadvantages: Wasteful to perform copy and then overwrite of memory

# Unix Process Creation

How are Unix shells implemented?

```
While (1) {
    Char *cmd = getcmd();
    Int retval = fork();
    If (retval == 0) {
        // This is the child process
        // Setup the child's process environment here
        // E.g., where is standard I/O, how to handle signals?
        exec(cmd);
        // exec does not return if it succeeds
        printf("ERROR: Could not execute %s\n", cmd);
        exit(1);
    } else {
        // This is the parent process; Wait for child to finish
        int pid = retval;
        wait(pid);
    }
}
```

# The fork() System Call

- Create a new process
  - The newly-created process has its own copy of the **address space, registers, and PC**.

p1.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {      // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {            // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

# Calling fork() example (Cont.)

Result (Not deterministic)

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

or

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

# The wait() System Call

- This system call won't return until the child has run and exited.

p2.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {      // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {            // parent goes down this path (main)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
               rc, wc, (int) getpid());
    }
    return 0;
}
```

# The wait() System Call (Cont.)

## Result (Deterministic)

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (wc:29267) (pid:29266)
prompt>
```

# The exec() System Call

- Run a program that is different from the calling program

p3.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {           // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {    // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc");           // program: "wc" (word count)
        myargs[1] = strdup("p3.c");         // argument: file to count
        myargs[2] = NULL;                 // marks end of array
        ...
    }
}
```

# The exec() System Call (Cont.)

## p3.c (Cont.)

```
...
    execvp(myargs[0], myargs); // runs word count
    printf("this shouldn't print out");
} else { // parent goes down this path (main)
    int wc = wait(NULL);
    printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
           rc, wc, (int) getpid());
}
return 0;
}
```

## Result

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
29 107 1030 p3.c
hello, I am parent of 29384 (wc:29384) (pid:29383)
prompt>
```

# All of the above with redirection

p4.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <fcntl.h>
#include <sys/wait.h>

int
main(int argc, char *argv[]){
    int rc = fork();
    if (rc < 0) {      // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child: redirect standard output to a file
        close(STDOUT_FILENO);
        open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
        ...
    }
}
```

# All of the above with redirection (Cont.)

p4.c

```
...
// now exec "wc"...
char *myargs[3];
myargs[0] = strdup("wc");           // program: "wc" (word count)
myargs[1] = strdup("p4.c");         // argument: file to count
myargs[2] = NULL;                  // marks end of array
execvp(myargs[0], myargs);        // runs word count
} else {                           // parent goes down this path (main)
    int wc = wait(NULL);
}
return 0;
}
```

## Result

```
prompt> ./p4
prompt> cat p4.output
32 109 846 p4.c
prompt>
```

# **Break**

# CPU Virtualization: Two Components


## Dispatcher

- Low-level mechanism
- Performs context-switch
  - Switch from user mode to kernel mode
  - Save execution state (registers) of old process in PCB
  - Insert PCB in ready queue
  - Load state of next process from PCB to registers
  - Switch from kernel to user mode
  - Jump to instruction in new user process

## Scheduler

- Policy to determine which process gets CPU when

# Review: State Transitions



How to transition? (“mechanism”)  
When to transition? (“policy”)

# Vocabulary

**Workload:** set of **job** descriptions (arrival time, run\_time)

- Job: Viewed as the current CPU burst of a process
- Process alternates between CPU and I/O  
process moves between ready and blocked queues

**Scheduler:** logic that decides which ready job to run

**Metric:** measurement of scheduling quality

# Scheduling Performance Metrics

## Minimize turnaround time

- Do not want to wait long for job to complete
- Completion\_time – arrival\_time

## Minimize response time

- Schedule interactive jobs promptly so users see output quickly
- Initial\_schedule\_time – arrival\_time

## Minimize waiting time

- Do not want to spend much time in Ready queue

## Maximize throughput

- Want many jobs to complete per unit of time

## Maximize resource utilization

- Keep expensive devices busy

## Minimize overhead

- Reduce number of context switches

## Maximize fairness

- All jobs get same amount of CPU over some time interval

# Workload Assumptions

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known

# Scheduling Basics

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time

# Example: workload, scheduler, metrics

JOB	arrival_time (s)	run_time (s)
A	~0	10
B	~0	10
C	~0	10

**FIFO:** First In, First Out

- also called FCFS (first come first served)
- run jobs in *arrival\_time* order

**What is our turnaround?:**  $completion\_time - arrival\_time$


# FIFO Event Trace

JOB	arrival_time (s)	run_time (s)	Time	Event
A	~0	10	0	A arrives
B	~0	10	0	B arrives
C	~0	10	0	C arrives
			10	run A
			10	complete A
			10	run B
			20	complete B
			20	run C
			30	complete C

# FIFO (Identical Jobs)

JOB	arrival_time (s)	run_time (s)
-----	------------------	--------------


A	~0	10
B	~0	10
C	~0	10



Gantt chart:

Illustrates how jobs are scheduled over time on a CPU


# FIFO (Identical Jobs)



What is the average turnaround time?

Def:  $\text{turnaround\_time} = \text{completion\_time} - \text{arrival\_time}$

# FIFO (Identical Jobs)



What is the average turnaround time?

Def:  $turnaround\_time = completion\_time - arrival\_time$   
 $(10 + 20 + 30) / 3 = 20s$

# Scheduling Basics

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time

# Workload Assumptions

1. ~~Each job runs for the same amount of time~~
2. All jobs arrive at the same time
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known

# Any Problematic Workloads for FIFO?

- **Workload:** ?
- **Scheduler:** FIFO
- **Metric:** turnaround is high


# Example: Big First Job

JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~0	10
C	~0	10

Draw Gantt chart for this workload and policy...

What is the average turnaround time?

# Example: Big First Job



Average turnaround time: **70s**

# Convoy Effect – Passing the Tractor



## Problem with Previous Scheduler:

FIFO: Turnaround time can suffer when short jobs must wait for long jobs

## New scheduler:

SJF (Shortest Job First)


Choose job with smallest *run\_time*

# Example: Shortest Job First

JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~0	10
C	~0	10

What is the average turnaround time with SJF?

# SJF Turnaround Time



What is the average turnaround time with SJF?

$$(80 + 10 + 20) / 3 = \text{\textcolor{red}{~36.7s}}$$

Average turnaround  
with FIFO: 70s

For minimizing average turnaround time (with no preemption): SJF is provably optimal

Moving shorter job before longer job improves turnaround time of short job more than it harms turnaround time of long job

# Scheduling Basics

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time

# Workload Assumptions


1. ~~Each job runs for the same amount of time~~
2. ~~All jobs arrive at the same time~~
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known

# Example: Shortest Job First (Arrival Time)

JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~10	10
C	~10	10

What is the average turnaround time with SJF?

# Stuck Behind that Tractor Again



JOB	arrival_time (s)	run_time (s)
A	~0	60
B	<b>~10</b>	10
C	<b>~10</b>	10

What is the average turnaround time?

$$(60 + (70 - 10) + (80 - 10)) / 3 = \mathbf{63.3s}$$

# Preemptive Scheduling


## Prev schedulers:

- FIFO and SJF are non-preemptive
- Only schedule new job when previous job voluntarily relinquishes CPU (performs I/O or exits)

## New scheduler:

- Preemptive: Potentially schedule different job at any point by taking CPU away from running job
- STCF (Shortest Time-to-Completion First)
- Always run a job that will complete the quickest

# Non-Preemptive: SJF




JOB	arrival_time (s)	run_time (s)
A	~0	60
B	<b>~10</b>	10
C	<b>~10</b>	10

What is the average turnaround time?

$$(60 + (70 - 10) + (80 - 10)) / 3 = \mathbf{63.3s}$$

# Preemptive: STCF



JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~10	10
C	~10	10

Average turnaround time with STCF?

**36.6s**

Average turnaround time with SJF: **63.3s**

# Scheduling Basics

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time


# Response Times

Sometimes care about when job starts instead of when it finishes

New metric:

$$\text{response\_time} = \text{first\_run\_time} - \text{arrival\_time}$$

# Response versus Turnaround



# Round Robin Scheduler


**Prev schedulers:**

FIFO, SJF, and STCF can have poor response time


**New scheduler:** RR (Round Robin)

Alternates ready processes every fixed-length time-slice

## FIFO versus RR



Avg Response Time?  
 $(0+5+10)/3 = 5$



Avg Response Time?  
 $(0+1+2)/3 = 1$

In what way is RR worse?

Average turn-around time with equal job lengths is horrible

Other reasons why RR could be better?

If don't know run-time of each job, gives short jobs a chance to run and finish fast

# Scheduling Basics

## Workloads:

arrival\_time  
run\_time

## Schedulers:

FIFO  
SJF  
STCF  
RR


## Metrics:

turnaround\_time  
response\_time

# Workload Assumptions


1. ~~Each job runs for the same amount of time~~
2. ~~All jobs arrive at the same time~~
3. ~~All jobs only use the CPU (no I/O)~~
4. Run-time of each job is known

# Not I/O Aware



Don't let Job A hold on to CPU while blocked waiting for disk

# I/O Aware (Overlap)



Treat Job A as 3 separate CPU bursts

When Job A completes I/O, another Job A is ready

Each CPU burst is shorter than Job B, so with SCTF,  
Job A preempts Job B

# Workload Assumptions

1. ~~Each job runs for the same amount of time~~
2. ~~All jobs arrive at the same time~~
3. ~~All jobs only use the CPU (no I/O)~~
4. ~~Run time of each job is known~~  
*(need smarter, fancier scheduler)*

# MLFQ (Multi-Level Feedback Queue)

Goal: general-purpose scheduling

Must support two job types with distinct goals


- “interactive” programs care about response time
- “batch” programs care about turnaround time

Approach: multiple levels of round-robin;  
each level has higher priority than lower levels  
and preempts them

# Priorities

Rule 1: If  $\text{priority}(A) > \text{Priority}(B)$ , A runs

Rule 2: If  $\text{priority}(A) == \text{Priority}(B)$ , A & B run in RR



“Multi-level”

How to know how to set priority?

Approach 1: nice

Approach 2: history “feedback”

# History

- Use past behavior of process to predict future behavior
  - Common technique in systems
- Processes alternate between **I/O** and **CPU** work
- Guess how CPU burst (job) will behave based on past CPU bursts (jobs) of this process

## More MLFQ Rules

Rule 1: If  $\text{priority}(A) > \text{Priority}(B)$ , A runs


Rule 2: If  $\text{priority}(A) == \text{Priority}(B)$ , A & B run in RR

More rules:


Rule 3: Processes start at top priority

Rule 4: If job uses whole slice, demote process  
(longer time slices at lower priorities)

# One Long Job Example




# An Interactive Process Joins



Interactive process never uses entire time slice, so never demoted


# Problems with MLFQ?



## Problems

- unforgiving + starvation
- gaming the system


# Prevent Starvation



Problem: Low priority job may never get scheduled

Periodically boost priority of all jobs (or all jobs that haven't been scheduled)

# Prevent Gaming



Problem: High priority job could trick scheduler and get more CPU by performing I/O right before time-slice ends

Fix: Account for job's total run time at priority level (instead of just this time slice); downgrade when exceeding threshold

# Proportional Share Scheduler – Lottery

- Fair-share scheduler
  - Guarantee that each job obtain a *certain percentage* of CPU time.
  - Not optimized for turnaround or response time

Approach:

- give processes lottery tickets
- whoever wins runs
- higher priority => more tickets

Amazingly simple to implement

# Basic Concept

- Tickets
  - Represent the share of a resource that a process should receive
  - The percent of tickets represents its share of the system resource in question.
- Example
  - There are two processes, A and B.
    - Process A has 75 tickets → receive 75% of the CPU
    - Process B has 25 tickets → receive 25% of the CPU

# Lottery scheduling

- The scheduler picks a winning ticket.
  - Load the state of that *winning process* and runs it.
- Example
  - There are 100 tickets
    - Process A has 75 tickets: 0 ~ 74
    - Process B has 25 tickets: 75 ~ 99

Scheduler's winning tickets: 63 85 70 39 76 17 29 41 36 39 10 99 68 83 63

Resulting scheduler: A B A A B A A A A A A B A B A

The longer these two jobs compete,  
The more likely they are to achieve the desired percentages.

# Ticket Mechanisms

- Ticket currency
  - A user allocates tickets among their own jobs in whatever currency they would like.
  - The system converts the currency into the correct global value.
- Example
  - There are 200 tickets (Global currency)
  - Process A has 100 tickets
  - Process B has 100 tickets

User A → 500 (A's currency) to A1 → 50 (global currency)  
→ 500 (A's currency) to A2 → 50 (global currency)


User B → 10 (B's currency) to B1 → 100 (global currency)

# Ticket Mechanisms (Cont.)

- Ticket transfer
  - A process can temporarily hand off its tickets to another process.
- Ticket inflation
  - A process can temporarily raise or lower the number of tickets it owns.
  - If any one process needs *more CPU time*, it can boost its tickets.

# Implementation

- Example: There are three processes, A, B, and C.
  - Keep the processes in a list:




```
1 // counter: used to track if we've found the winner yet
2 int counter = 0;
3
4 // winner: use some call to a random number generator to
5 // get a value, between 0 and the total # of tickets
6 int winner = getrandom(0, totaltickets);
7
8 // current: use this to walk through the list of jobs
9 node_t *current = head;
10
11 // loop until the sum of ticket values is > the winner
12 while (current) {
13     counter = counter + current->tickets;
14     if (counter > winner)
15         break; // found the winner
16     current = current->next;
17 }
18 // 'current' is the winner: schedule it...
```

# Implementation (Cont.)

- $U$ : unfairness metric
  - The time the first job completes divided by the time that the second job completes.
- Example:
  - There are two jobs, each job has runtime 10.
  - First job finishes at time 10
  - Second job finishes at time 20
  - $U = \frac{10}{20} = 0.5$
  - $U$  will be close to 1 when both jobs finish at nearly the same time.

# Lottery Fairness Study

- There are two jobs.
- Each job has the same number of tickets (100).



When the job length is not very long,  
average unfairness can be quite severe.

# Stride Scheduling

- Stride of each process
  - (A large number) / (the number of tickets of the process)
  - Example: A large number = 10,000
    - Process A has 100 tickets → stride of A is 100
    - Process B has 50 tickets → stride of B is 200
- A process runs, increment a counter(=pass value) for it by its stride.
- Pick the process to run that has the lowest pass value

```
current = remove_min(queue);      // pick client with minimum pass
schedule(current);              // use resource for quantum
current->pass += current->stride; // compute next pass using stride
insert(queue, current);         // put back into the queue
```

A pseudo code implementation

# Summary

- Understand goals (metrics) and workload, then design scheduler around that
- General purpose schedulers need to support processes with different goals
- Past behavior is good predictor of future behavior
- Random algorithms (lottery scheduling) can be simple to implement and avoid corner cases.

# What's Next

- Module Software Preparation (HW0)
  - Due by Sunday, January 25 by 11:59pm ET
  - Install Virtual Machine ([VirtualBox or UTM or something else](#))
  - Install GitHub repository
  - Verify C Compiler works
  - Push file into GitHub private repo
- Programming Assignment 1
  - Due by Sunday, February 1 by 11:59pm ET
  - Write a [UNIX Shell](#)
- Homework Assignment 1
  - Due by Sunday, February 1 by 11:59pm ET
- Quiz 1
  - Due by Sunday, February 1 by 11:59pm ET