

CS 5600 Computer Systems

Spring 2026

Virtualization: Memory

Lecture 4.1

February 3, 2026

Prof. Scott Valcourt
s.valcourt@northeastern.edu
603-380-2860 (cell)

Portland, ME



Memory Virtualization


- What is **memory virtualization**?
 - OS virtualizes its physical memory.
 - OS provides an **illusion memory space** per each process.
 - It seems like **each process uses the whole memory**.

Benefits of Memory Virtualization

- Ease of use in programming
- Memory efficiency in terms of **time** and **space**
- The guarantee of isolation for processes as well as OS
 - Protection from **errant accesses** of other processes


OS in the Early System Years

- Load only one process in memory.
 - Poor utilization and efficiency




Multiprogramming and Time Sharing

- Load multiple processes in memory.
 - Execute one for a short while.
 - Switch processes between them in memory.
 - Increase utilization and efficiency.
- Causes an important protection issue.
 - Errant memory accesses from other processes




Address Space

- OS creates an **abstraction** of physical memory.
 - Address space contains everything about a running process.
 - That is program code, heap, stack, etc.



Address Space (Cont.)

- Code
 - Where instructions live
- Heap
 - Dynamically allocate memory.
 - `malloc` in C language
 - `new` in object-oriented language
- Stack
 - Store return addresses or values.
 - Contain local variables and arguments to routines.



Virtual Address

- **Every address** in a running program is virtual.
 - OS translates the virtual address to physical address

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]){

    printf("location of code : %p\n", (void *) main);
    printf("location of heap : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack : %p\n", (void *) &x);


    return x;
}
```

A simple program that prints out addresses

Virtual Address (Cont.)

- The output in 64-bit Linux machine

```
location of code : 0x40057d
location of heap : 0xcf2010
location of stack : 0x7fff9ca45fcc
```



Memory API: malloc()

```
#include <stdlib.h>

void* malloc(size_t size)
```

- Allocate a memory region on the heap.
 - Argument
 - `size_t size`: size of the memory block(in bytes)
 - `size_t` is an unsigned integer type.
 - Return
 - Success: a void type pointer to the memory block allocated by `malloc`
 - Fail: a NULL pointer

Memory API: sizeof()

- Routines and macros are utilized for size in malloc instead typing in a number directly.
- Two types of results of sizeof with variables
 - The actual size of 'x' is known at run-time.

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

4

- The actual size of 'x' is known at compile-time.

```
int x[10];  
printf("%d\n", sizeof(x));
```

40



Memory API: free()

```
#include <stdlib.h>


void free(void* ptr)
```

- Free a memory region allocated by a call to malloc.
 - Argument
 - void *ptr : a pointer to a memory block allocated with malloc
 - Return
 - none

Memory Allocating




Memory Freeing



Forgetting To Allocate Memory

- Incorrect code


```
char *src = "hello"; //character string constant
char *dst;           //unallocated
strcpy(dst, src);   //segfault and die
```



Forgetting To Allocate Memory(Cont.)

- Correct code


```
char *src = "hello"; //character string constant  
char *dst (char *)malloc(strlen(src) + 1 ); // allocated  
strcpy(dst, src); //work properly
```



Not Allocating Enough Memory

- Incorrect code, but works properly


```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src)); // too small
strcpy(dst, src); //work properly
```



Forgetting to Initialize


- Encounter an uninitialized read


```
int *x = (int *)malloc(sizeof(int)); // allocated  
printf("*x = %d\n", *x);           // uninitialized memory access
```



Memory Leak




- A program runs out of memory and eventually dies.

 unused : unused, but not freed




Dangling Pointer

- Freeing memory before it is finished using
 - A program accesses memory with an invalid pointer



`free(b)`



Other Memory APIs: calloc()

```
#include <stdlib.h>


void *calloc(size_t num, size_t size)
```

- Allocate memory on the heap and zeroes it before returning.
 - Argument
 - `size_t num` : number of blocks to allocate
 - `size_t size` : size of each block(in bytes)
 - Return
 - Success: a void type pointer to the memory block allocated by `calloc`
 - Fail: a `NULL` pointer

Double Free

- Free memory that was freed already.

```
int *x = (int *)malloc(sizeof(int)); // allocated
free(x);                            // free memory
free(x);                            // free repeatedly
```



Other Memory APIs: realloc()

```
#include <stdlib.h>

void *realloc(void *ptr, size_t size)
```

- Change the size of memory block.
 - A pointer returned by `realloc` may be either the same as `ptr` or a new one.
 - Argument
 - `void *ptr`: Pointer to memory block allocated with `malloc`, `calloc`, or `realloc`
 - `size_t size`: New size for the memory block (in bytes)
 - Return
 - Success: Void type pointer to the memory block
 - Fail: a `NULL` pointer

System Calls

```
#include <unistd.h>

int brk(void *addr)
void *sbrk(intptr_t increment);
```

- malloc library call use `brk` system call.
 - `brk` is called to expand the program's *break*.
 - *break*: The location of **the end of the heap** in address space
 - `sbrk` is an additional call similar to `brk`.
 - Programmers **should never directly call** either `brk` or `sbrk`.

System Calls (Cont.)

```
#include <sys/mman.h>

void *mmap(void *ptr, size_t length, int port, int flags,
int fd, off_t offset)
```

- mmap system call can create **an anonymous** memory region.

Quiz Time! Match that Address Location

```
int x;  
int main(int argc, char *argv[] ) {  
    int y;  
    int *z = malloc(sizeof(int));}  
}
```

Possible segments: static data, code, stack, heap

Address	Location
x	Static data → Code
main	Code
y	Stack
z	Stack
*z	Heap

What if no static data segment exists?

Break

Lecture Agenda

- Chapter 15 – Address Translation
- Chapter 16 – Segmentation
- Chapter 17 – Free Space Management
- Chapter 18 – Paging
- Chapter 19 – Translation Lookaside Buffers
- Chapter 20 – Paging: Smaller Tables
- Chapter 21 – Swapping: Mechanisms
- Chapter 22 – Swaping: Policies

Memory Virtualizing with Efficiency and Control

- Memory virtualizing takes a similar strategy known as **limited direct execution(LDE)** for efficiency and control.
- In memory virtualizing, efficiency and control are attained by **hardware support**.
 - e.g., registers, TLB(Translation Look-aside Buffer)s, page-table

Address Translation

- Hardware transforms a **virtual address** to a **physical address**.
 - The desired information is actually stored in a physical address.
- The OS must get involved at key points to set up the hardware.
 - The OS must manage memory to judiciously intervene.

Example: Address Translation

- C - Language code

```
void func()
    int x;
    ...
    x = x + 3;      // this is the line of code we are interested in
```

- **Load** a value from memory
- **Increment** it by three
- **Store** the value back into memory


Example: Address Translation (Cont.)

- Assembly

```
128 : movl 0x0(%ebx), %eax      ; load 0+ebx into eax
132 : addl $0x03, %eax          ; add 3 to eax register
135 : movl %eax, 0x0(%ebx)      ; store eax back to mem
```

- Presume that the address of 'x' has been placed in ebx register.
- Load the value at that address into eax register.
- Add 3 to eax register.
- Store the value in eax back into memory.

Example: Address Translation (Cont.)




- Fetch instruction at address 128
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)


Relocation Address Space

- The OS wants to place the process **somewhere else** in physical memory, not at address 0.
 - The address space start at address 0.

A Single Relocated Process



Base and Bounds Register



Dynamic (Hardware Base) Relocation

- When a program starts running, the OS decides **where** in physical memory a process should be **loaded**.
- Set the **base** register to a value.

$$\text{physical address} = \text{virtual address} + \text{base}$$

- Every virtual address must **not be greater than the bound** and **negative**.

$$0 \leq \text{virtual address}$$
$$\text{virtual address} < \text{bounds}$$

Relocation and Address Translation


```
128 : movl 0x0(%ebx), %eax
```

- Fetch instruction at address 128


$$32896 = 128 + 32KB(base)$$

- Execute this instruction
 - Load from address 15KB

$$47KB = 15KB + 32KB(base)$$



Two Ways of Bounds Register




OS Issues for Memory Virtualizing

- The OS must **take action** to implement **base-and-bounds** approach.
- Three critical junctures:
 - When a process **starts running**:
 - Finding space for address space in physical memory
 - When a process is **terminated**:
 - Reclaiming the memory for use
 - When context **switch occurs**:
 - Saving and storing the base-and-bounds pair


OS Issues: When a Process Starts Running

- The OS must **find a room** for a new address space.
 - free list : A list of the range of the physical memory which are not in use.




OS Issues: When a Process Is Terminated

- The OS must put the memory back on the free list.




OS Issues: When Context Switch Occurs

- The OS must **save and restore** the base-and-bounds pair.
 - In **process structure or process control block(PCB)**



Inefficiency of the Base and Bound Approach




- Big chunk of “free” space
- “free” space takes up physical memory.
- Hard to run when an address space does not fit into physical memory

Segmentation

- Segment is just a **contiguous portion** of the address space of a particular length.
 - Logically-different segment: code, stack, heap
- Each segment can be **placed in different part of physical memory.**
 - **Base and bounds exist per each segment.**

Placing Segment In Physical Memory




Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Address Translation on Segmentation

$$\text{physical address} = \text{offset} + \text{base}$$


- The offset of virtual address 100 is 100.
 - The code segment **starts at virtual address 0** in address space.



Address Translation on Segmentation (Cont.)


Virtual address + base is not the correct physical address.

- The offset of virtual address 4200 is 104.
 - The heap segment **starts at virtual address 4096** in address space.



Segmentation Fault or Violation


- If an **illegal address** such as 7KB which is beyond the end of heap is referenced, the OS occurs a **segmentation fault**.
 - The hardware detects that address is **out of bounds**.



Referring to Segment


- **Explicit approach**

- Chop up the address space into segments based on the **top few bits** of the virtual address.



- Example: virtual address 4200 (01000001101000)

Segment	bits
Code	00
Heap	01
Stack	10
-	11




Referring to Segment (Cont.)

```
1 // get top 2 bits of 14-bit VA
2 Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3 // now get offset
4 Offset = VirtualAddress & OFFSET_MASK
5 if (Offset >= Bounds[Segment])
6 RaiseException(PROTECTION_FAULT)
7 else
8 PhysAddr = Base[Segment] + Offset
9 Register = AccessMemory(PhysAddr)
```

- SEG_MASK = 0x3000 (1100000000000000)
- SEG_SHIFT = 12
- OFFSET_MASK = 0xFF (001111111111)

Referring to Stack Segment

- Stack grows **backward**.
- **Extra hardware support** is need.
 - The hardware checks which way the segment grows.
 - 1: positive direction, 0: negative direction



Support for Sharing

- Segment can be **shared between address space**.
 - **Code sharing** is still in use in systems today.
 - by extra hardware support.
- Extra hardware support is need for form of **Protection bits**.
 - **A few more bits** per segment to indicate **permissions** of **read, write and execute**.

Segment Register Values (with Protection)

Segment	Base	Size	Grows	Positive?	Protection
Code	32K	2K		1	Read-Execute
Heap	34K	2K		1	Read-Write
Stack	28K	2K		0	Read-Write


Fine-Grained and Coarse-Grained

- **Coarse-Grained** means segmentation in a small number.
 - e.g., code, heap, stack.
- **Fine-Grained** segmentation allows **more flexibility** for the address space in some early systems.
 - To support many segments, hardware support with a **segment table** is required.

OS Support: Fragmentation


- **External Fragmentation:** little holes of **free space** in physical memory that make it difficult to allocate new segments.
 - There is **24KB free**, but **not in one contiguous** segment.
 - The OS **cannot** satisfy the **20KB request**.
- **Compaction:** rearranging the exiting segments in physical memory.
 - Compaction is **costly**.
 - Stop running process.
 - Copy data to somewhere.
 - Change segment register value.

Memory Compaction





Splitting

- Finding a free chunk of memory that can satisfy the request and splitting it into two.
 - When request for memory allocation is **smaller** than the size of free chunks.




Splitting (Cont.)

- Two 10-bytes free segment with 1-byte request



Coalescing


- If a user requests memory that is **bigger than free chunk size**, the list will **not find** such a free chunk.
- Coalescing: **Merge** returning a free chunk with existing chunks into a large single free chunk if **addresses** of them are **nearby**.



Tracking The Size of Allocated Regions


- The interface to `free (void *ptr)` does **not take a size parameter**.
 - How does the library **know the size** of the memory region that will be back **into the free list**?
- Most allocators store **extra information** in a **header** block.

```
ptr = malloc(20);
```



The Header of Allocated Memory Chunk

- The header minimally **contains the size** of the allocated memory region.
- The header may also contain:
 - Additional pointers to speed up de-allocation
 - A magic number for integrity checking



```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

A Simple Header

The Header of Allocated Memory Chunk (Cont.)

- The **size** for free region is the **size of the header plus the size of the space** allocated to the user.
 - If a user **requests N bytes**, the library searches for a free chunk of **size N plus the size of the header**
 - Simple pointer arithmetic to find the header pointer.

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```

Embedding A Free List

- The memory-allocation library **initializes** the heap and **puts** the first element of **the free list** in **the free space**.
 - The library **can't use** `malloc()` to build a list **within itself**.

Embedding A Free List (Cont.)

- Description of a node of the list


```
typedef struct __node_t {
    int size;
    struct __node_t *next;
} nodet_t;
```

- Building heap and putting a free list
 - Assume that the heap is built via the `mmap()` system call.

```
// mmap() returns a pointer to a chunk of free space
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,
                    MAP_ANON|MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```

A Heap With One Free Chunk

```
// mmap() returns a pointer to a chunk of free space
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,
                    MAP_ANON|MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```




Embedding A Free List: Allocation


- If a chunk of memory is requested, the library **will first find** a chunk that is **large enough** to accommodate the request.
- The library will
 - **Split** the large free chunk into two.
 - One for the **request** and the **remaining** free chunk
 - **Shrink** the size of the free chunk in the list.

Embedding A Free List: Allocation (Cont.)

- Example: a request for 100 bytes by `ptr = malloc(100)`
 - Allocating 108 bytes out of the existing one free chunk.
 - Shrinking the one free chunk to 3980 (4088 minus 108).



Free Space With Chunks Allocated




Free Space With Three Chunks Allocated

Free Space With `free()`


Example: `free(sptr)`

- The 100 bytes chunk is **back into** the free list.
- The free list will **start** with a **small chunk**.
 - The list header will point to the small chunk




Free Space With Freed Chunks

- Let's assume that the last two in-use chunks are freed.
- **External Fragmentation** occurs.
 - Coalescing is needed in the list.



Growing The Heap

- Most allocators **start with a small-sized heap** and then **request more** memory from the OS when they run out.
 - e.g., `sbrk()`, `brk()` in most UNIX systems.



Managing Free Space: Basic Strategies


- Best Fit:
 - Finding free chunks that are **big or bigger than the request**
 - Returning the **one of smallest** in the chunks **in the group** of candidates
- Worst Fit:
 - Finding the **largest free chunks** and allocation the amount of the request
 - **Keeping the remaining chunk** on the free list.

Managing Free Space: Basic Strategies (Cont.)


- First Fit:
 - Finding the **first chunk** that is **big enough** for the request
 - Returning the requested amount and remaining the rest of the chunk.
- Next Fit:
 - Finding the first chunk that is big enough for the request.
 - Searching at **where one was looking** at instead of the beginning of the list.

Examples of Basic Strategies


- Allocation Request Size 15



- Result of Best-fit



- Result of Worst-fit



Other Approaches: Segregated List


- Segregated List:
 - Keeping free chunks in different size in a separate list for the size of popular request.
 - New Complication:
 - **How much** memory should dedicate to **the pool of memory** that serves **specialized requests** of a given size?
 - **Slab allocator** handles this issue.

Other Approaches: Segregated List (Cont.)

- Slab Allocator
 - Allocate several object caches.
 - The objects are likely to be requested frequently.
 - e.g., locks, file-system inodes, etc.
 - **Request some memory from a more general memory allocator when a given cache is running low** on free space.

Other Approaches: Buddy Allocation

- Binary Buddy Allocation
 - The allocator **divides free space by two until a block** that is big enough to accommodate the request is **found**.



Other Approaches: Buddy Allocation(Cont.)

- Buddy allocation can suffer from **internal fragmentation**.
- Buddy system makes **coalescing** simple.
 - **Coalescing** two blocks into the next level of the block.

Break

Concept of Paging



- Paging **splits up** address space into **fixed-sized** unit called a **page**.
 - Segmentation: variable size of logical segments(code, stack, heap, etc.)
- With paging, **physical memory** is also **split** into some number of pages called a **page frame**.
- **Page table** per process is needed **to translate** the virtual address to physical address.

Advantages of Paging

- **Flexibility:** Supporting the abstraction of address space effectively
 - Don't need assumption how heap and stack grow and are used.
- **Simplicity:** ease of free-space management
 - The page in address space and the page frame are the same size.
 - Easy to allocate and keep a free list

Example: A Simple Paging


- 128-byte physical memory with 16-byte page frames
- 64-byte address space with 16-byte pages




64-Byte Address Space Placed In Physical Memory

Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page




- Example: virtual address 21 in 64-byte address space



Example: Address Translation


- The virtual address 21 in 64-byte address space



Where are Page Tables Stored?

- Page tables can get awfully large
 - 32-bit address space with 4-KB pages, 20 bits for VPN
 - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- Page tables for each process are stored in memory.

Example: Page Table in Kernel Physical Memory




What is in the Page Table?

- The page table is just a **data structure** that is used to map the virtual address to physical address.
 - Simplest form: a linear page table, an array
- The OS **indexes** the array by VPN and looks up the page-table entry.

Common Flags of Page Table Entries

- **Valid Bit:** Indicating whether the particular translation is valid
- **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- **Present Bit:** Indicating whether this page is in physical memory or on disk (swapped out)
- **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- **Reference Bit (Accessed Bit):** Indicating that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- P: present
- R/W: read/write bit
- U/S: supervisor
- A: accessed bit
- D: dirty bit
- PFN: the page frame number

Paging: Too Slow

- To find a location of the desired PTE, the **starting location** of the page table is **needed**.
- For every memory reference, paging requires the OS to perform one **extra memory reference**.

Accessing Memory With Paging

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

A Memory Trace

- Example: A Simple Memory Access

```
int array[1000];
...
for (i = 0; i < 1000; i++)
    array[i] = 0;
```


- Compile and execute

```
prompt> gcc -o array array.c -Wall -o
prompt>./array
```

- Resulting Assembly code


```
0x1024 movl $0x0, (%edi,%eax,4)
0x1028 incl %eax
0x102c cmpl $0x03e8,%eax
0x1030 jne 0x1024
```

A Virtual (and Physical) Memory Trace



Translation Lookaside Buffer (TLB)

- Part of the chip's memory-management unit(MMU).
- A hardware cache of **popular** virtual-to-physical address translation.



TLB Basic Algorithms

```
1: VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2: (Success, TlbEntry) = TLB_Lookup(VPN)
3: if (Success == True) { // TLB Hit
4:   if (CanAccess(TlbEntry.ProtectBit) == True) {
5:     offset = VirtualAddress & OFFSET_MASK
6:     PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:     AccessMemory(PhysAddr)
8:   } else RaiseException(PROTECTION_ERROR)
```

- (line 1) extract the virtual page number (VPN).
- (line 2) check if the TLB holds the translation for this VPN.
- (lines 5-8) extract the page frame number from the relevant TLB entry, and from the desired physical address and access memory.

TLB Basic Algorithms (Cont.)

```
11:     }else{ //TLB Miss  
12:         PTEAddr = PTBR + (VPN * sizeof(PTE))  
13:         PTE = AccessMemory(PTEAddr)  
14:         (...)  
15:     }else{  
16:         TLB_Insert( VPN , PTE.PFN , PTE.ProtectBits)  
17:         RetryInstruction()  
18:     }  
19: }
```

- (lines 11-12) The hardware accesses the page table to find the translation.
- (line 16) updates the TLB with the translation.

Example: Accessing An Array

- How a TLB can improve its performance

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					


```
0: int sum = 0 ;
1: for( i=0; i<10; i++) {
2:     sum+=a[i];
3: }
```

The TLB improves performance
due to spatial locality


3 misses and 7 hits.
Thus TLB hit rate is 70%.

Locality

- Temporal Locality
 - An instruction or data item that has been recently accessed will likely be re-accessed soon in the future.



- Spatial Locality
 - If a program accesses memory at address x, it will likely soon access memory near x.



Who Handles the TLB Miss?

- Hardware handle the TLB miss entirely on CISC.
 - The hardware has to know exactly where the page tables are located in memory.
 - The hardware would “walk” the page table, find the correct page-table entry and extract the desired translation, update and retry instruction.
 - hardware-managed TLB.

Who Handles the TLB Miss? (Cont.)

- RISC have what is known as a **software-managed TLB**.
 - On a TLB miss, the hardware raises exception(trap handler).
 - **Trap handler is code** within the OS that is written with the express purpose of handling TLB miss.

TLB Control Flow Algorithm (OS Handled)

```
1:     VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2:     (Success, TlbEntry) = TLB_Lookup(VPN)
3:     if (Success == True) // TLB Hit
4:         if (CanAccess(TlbEntry.ProtectBits) == True)
5:             Offset = VirtualAddress & OFFSET_MASK
6:             PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:             Register = AccessMemory(PhysAddr)
8:         else
9:             RaiseException(PRETECTION_FAULT)
10:    else // TLB Miss
11:        RaiseException(TLB_MISS)
```


TLB Entry

- TLB is managed by a **Full Associative** method.
 - A typical TLB might have 32, 64, or 128 entries.
 - Hardware search the entire TLB in parallel to find the desired translation.
 - other bits: valid bits, protection bits, address-space identifier, dirty bit




Typical TLB entry look like this

TLB Issue: Context Switching





TLB Issue: Context Switching



TLB Issue: Context Switching

Process A




TLB Table

VPN	PFN	valid	prot
10	100	1	rwx
-	-	-	-
10	170	1	rwx
-	-	-	-

Can't distinguish which entry is meant for which process

To Solve the Problem

- Provide an address space identifier(ASID) field in the TLB.



Another Case


- Two processes **share a page**.
 - Process 1 is sharing physical page 101 with Process2.
 - P1 maps this page into the 10th page of its address space.
 - P2 maps this page to the 50th page of its address space.

VPN	PFN	valid	prot	ASID
10	101	1	rwx	1
-	-	-	-	-
50	101	1	rwx	2
-	-	-	-	-

Sharing of pages is **useful** as it reduces the number of physical pages in use.


TLB Replacement Policy

- LRU (Least Recently Used)
 - Evict an entry that has not recently been used.
 - Take advantage of *locality* in the memory-reference stream.



A Real TLB Entry


All 64 bits of this TLB entry(example of MIPS R4000)



Flag	Content
19-bit VPN	The rest reserved for the kernel.
24-bit PFN	Systems can support with up to 64GB of main memory($2^{24} * 4KB$ pages).
Global bit(G)	Used for pages that are globally-shared among processes.
ASID	OS can use to distinguish between address spaces.
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.

Paging: Linear Tables

- We usually have one page table for every process in the system.
- Assume that 32-bit address space with 4KB pages and 4-byte page-table entry.




$$\text{Page table size} = \frac{2^{32}}{2^{12}} * 4\text{Byte} = 4\text{MByte}$$

Page tables are **too big** and thus consume too much memory.

Paging: Smaller Tables

- Page tables are too big and thus consume too much memory.
- Assume that 32-bit address space with **16KB** pages and 4-byte page-table entry.




$$\frac{2^{32}}{2^{16}} * 4 = 1MB \text{ per page table}$$

Big pages lead to internal fragmentation.


Problem

- Single page table for the entries address space of process.



Problem

- Most of the page table is **unused**, full of invalid entries.




Hybrid Approach: Paging and Segments

- In order to reduce the memory overhead of page tables.
 - Using base not to point to the segment itself but rather to hold the **physical address of the page table** of that segment.
 - The bounds register is used to indicate the end of the page table.

Simple Example of Hybrid Approach

- Each process has **three** page tables associated with it.
 - When process is running, the base register for each of these segments contains the physical address of a linear page table for that segment.



Seg value	Content
00	unused segment
01	code
10	heap
11	stack

TLB miss on Hybrid Approach

- The hardware get to **physical address** from **page table**.
 - The hardware uses the segment bits(SN) to determine which base and bounds pair to use.
 - The hardware then takes the **physical address** therein and **combines** it with the VPN as follows to form the address of the page table entry(PTE) .

```
01: SN = (VirtualAddress & SEG_MASK) >> SN_SHIFT  
02: VPN = (VirtualAddress & VPN_MASK) >> VPN_SHIFT  
03: AddressOfPTE = Base[SN] + (VPN * sizeof(PTE))
```


Problem of Hybrid Approach

- Hybrid Approach is not without problems.
 - If we have a large but sparsely-used heap, we can still end up with a lot of page table waste.
 - Causing external fragmentation to arise again.

Multi-level Page Tables

- Turns the linear page table into something like a tree.
 - Chop up the page table into page-sized units.
 - If an entire page of page-table entries is invalid, don't allocate that page of the page table at all.
 - To track whether a page of the page table is valid, use a new structure, called **page directory**.

Multi-level Page Tables: Page Directory



Multi-level Page Tables: Page Directory Entries

- The page directory contains one entry per page of the page table.
 - It consists of a number of **page directory entries (PDE)**.
- PDE has a valid bit and page frame number (PFN).

Multi-level Page Tables: Advantages & Disadvantages



- Advantage
 - Only allocates page-table space in proportion to the amount of address space you are using.
 - The OS can grab the next free page when it needs to allocate or grow a page table.
- Disadvantage
 - Multi-level table is a small example of a time-space trade-off.
 - Complexity.

Multi-level Page Table: Level of Indirection

- A multi-level structure can adjust **level of indirection** through use of the page directory.
 - Indirection place page-table pages wherever we would like in physical memory.


A Detailed Multi-Level Example

- To understand the idea behind multi-level page tables better, let's do an example.




A Detailed Multi-Level Example: Page Directory Index

- The page directory needs one entry per page of the page table
 - it has 16 entries.
- The page-directory entry is invalid → Raise an exception (The access is invalid)




A Detailed Multi-Level Example: Page Table Index

- The PDE is valid, we have more work to do.
 - To fetch the page table entry (PTE) from the page of the page table pointed to by this page-directory entry.
- This **page-table index** can then be used to index into the page table itself.



More than Two Levels


- In some cases, a deeper tree is possible.



Flag	Detail
Virtual address	30 bit
Page size	512 byte
VPN	21 bit
Offset	9 bit

More than Two Levels : Page Table Index

- In some cases, a deeper tree is possible.




Flag	Detail
Virtual address	30 bit
Page size	512 byte
VPN	21 bit
Offset	9 bit
Page entry per page	128 PTEs

$$\log_2 128 = 7$$

More than Two Levels: Page Directory

- If our page directory has 2^{14} entries, it spans not one page but 128.
- To remedy this problem, we build a **further level** of the tree, by splitting the page directory itself into multiple pages of the page directory.



Multi-level Page Table Control Flow

```
01: VPN = (VirtualAddress & VPN_MASK) >> SHIFT
02: (Success,TlbEntry) = TLB_Lookup(VPN)
03: if(Success == True) //TLB Hit
04:   if(CanAccess(TlbEntry.ProtectBits) == True)
05:     Offset = VirtualAddress & OFFSET_MASK
06:     PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
07:     Register = AccessMemory(PhysAddr)
08:   else RaiseException(PROTECTION_FAULT);
09: else // perform the full multi-level lookup
```

- ◆ (line 1) extract the virtual page number (VPN)
- ◆ (line 2) check if the TLB holds the translation for this VPN
- ◆ (lines 5-8) extract the page frame number from the relevant TLB entry, and form the desired physical address and access memory

Multi-level Page Table Control Flow

```
11: else
12:     PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13:     PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14:     PDE = AccessMemory(PDEAddr)
15:     if(PDE.Valid == False)
16:         RaiseException(SEGMENTATION_FAULT)
17:     else // PDE is Valid: now fetch PTE from PT
```

- (line 11) extract the Page Directory Index (PDIndex)
- (line 13) get Page Directory Entry (PDE)
- (lines 15-17) Check PDE valid flag. If valid flag is true, fetch Page Table entry from Page Table

The Translation Process: Remember the TLB


```
18: PTIndex = (VPN & PT_MASK) >> PT_SHIFT  
19: PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))  
20: PTE = AccessMemory(PTEAddr)  
21: if(PTE.Valid == False)  
22:     RaiseException(SEGMENTATION_FAULT)  
23: else if(CanAccess(PTE.ProtectBits) == False)  
24:     RaiseException(PROTECTION_FAULT);  
25: else  
26:     TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)  
27:     RetryInstruction()
```

Inverted Page Tables

- Keeping a single page table that has an entry for each physical page of the system.
- The entry tells us which process is using this page, and which virtual page of that process maps to this physical page.

Beyond Physical Memory: Mechanisms

- Require an additional level in the **memory hierarchy**.
 - OS need a place to stash away portions of address space that currently aren't in great demand.
 - In modern systems, this role is usually served by a **hard disk drive**




Single large address for a process

- Always need to first arrange for the code or data to be in memory when before calling a function or accessing data.
- To Beyond just a single process.
 - The addition of swap space allows the OS to support the illusion of a large virtual memory for multiple concurrently-running process

Swap Space

- Reserve some space on the disk for moving pages back and forth.
- OS need to remember to the swap space, in page-sized unit



Present Bit

- Add some machinery higher up in the system in order to support swapping pages to and from the disk.
 - When the hardware looks in the PTE, it may find that the page is not present in physical memory.

Value	Meaning
1	page is present in physical memory
0	The page is not in memory but rather on disk.

What If Memory Is Full?


- The OS likes to page out pages to make room for the new pages the OS is about to bring in.
 - The process of picking a page to kick out, or replace is known as **page-replacement** policy

The Page Fault

- Accessing a page that is **not in physical memory**.
 - If a page is not present and has been swapped disk, the OS need to swap the page into memory in order to service the page fault.

Page Fault Control Flow

- PTE used for data such as the PFN of the page for a disk address.



Page Fault Control Flow – Hardware

```
1:   VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2:   (Success, TlbEntry) = TLB_Lookup(VPN)
3:   if (Success == True) // TLB Hit
4:     if (CanAccess(TlbEntry.ProtectBits) == True)
5:       Offset = VirtualAddress & OFFSET_MASK
6:       PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:       Register = AccessMemory(PhysAddr)
8:     else RaiseException(PRODUCTION_FAULT)
```

Page Fault Control Flow – Hardware

```
9: else // TLB Miss
10: PTEAddr = PTBR + (VPN * sizeof(PTE))
11: PTE = AccessMemory(PTEAddr)
12: if (PTE.Valid == False)
13:     RaiseException(SEGMENTATION_FAULT)
14: else
15: if (CanAccess(PTE.ProtectBits) == False)
16:     RaiseException(PROTECTION_FAULT)
17: else if (PTE.Present == True)
18: // assuming hardware-managed TLB
19:     TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
20:     RetryInstruction()
21: else if (PTE.Present == False)
22:     RaiseException(PAGE_FAULT)
```

Page Fault Control Flow – Software

```
1:     PFN = FindFreePhysicalPage()
2:     if (PFN == -1) // no free page found
3:         PFN = EvictPage() // run replacement algorithm
4:         DiskRead(PTE.DiskAddr, pfn) // sleep (waiting for I/O)
5:         PTE.present = True // update page table with present
6:         PTE.PFN = PFN // bit and translation (PFN)
7:         RetryInstruction() // retry instruction
```

- ◆ The OS must find a physical frame for the soon-be-faulted-in page to reside within.
- ◆ If there is no such page, waiting for the replacement algorithm to run and kick some pages out of memory.

When Replacements Really Occur

- OS waits until memory is entirely full, and only then replaces a page to make room for some other page
 - This is a little bit unrealistic, and there are many reason for the OS to keep a small portion of memory free more proactively.
- Swap Daemon, Page Daemon
 - There are fewer than LW pages available, a background thread that is responsible for freeing memory runs.
 - The thread evicts pages until there are HW pages available.

Beyond Physical Memory: Policies

- Memory pressure forces the OS to start **paging out** pages to make room for actively-used pages.
- Deciding which page to evict is encapsulated within the replacement policy of the OS.

Cache Management

- Goal in picking a replacement policy for this cache is to minimize the number of cache misses.
- The number of cache hits and misses let us calculate the *average memory access time(AMAT)*.

$$AMAT = (P_{Hit} * T_M) + (P_{Miss} * T_D)$$

Argument	Meaning
T_M	The cost of accessing memory
T_D	The cost of accessing disk
P_{Hit}	The probability of finding the data item in the cache(a hit)
P_{Miss}	The probability of not finding the data in the cache(a miss)

The Optimal Replacement Policy

- Leads to the fewest number of misses overall
 - Replaces the page that will be accessed furthest in the future
 - Resulting in the **fewest-possible** cache misses
- Serve only as a comparison point, to know how close we are to **perfect**

Tracing the Optimal Policy

Reference Row										
0	1	2	0	1	3	0	3	1	2	1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	2	0,1,3
0	Hit		0,1,3
3	Hit		0,1,3
1	Hit		0,1,3
2	Miss	3	0,1,2
1	Hit		0,1,2

$$\text{Hit rate is } \frac{\text{Hits}}{\text{Hits}+\text{Misses}} = 54.6\%$$

Future is not known.

A Simple Policy: FIFO

- Pages were placed in a queue when they enter the system.
- When a replacement occurs, the page on the tail of the queue(the “**First-in**” pages) is evicted.
 - It is simple to implement but can't determine the importance of blocks.

Tracing the FIFO Policy

Reference Row										
0	1	2	0	1	3	0	3	1	2	1


Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss		3,0,1
2	Miss	3	0,1,2
1	Hit		0,1,2

$$\text{Hit rate is } \frac{\text{Hits}}{\text{Hits+Misses}} = 36.4\%$$

Even though page 0 had been accessed a number of times, FIFO still kicks it out.

BELADY'S ANOMALY

- We would expect the cache hit rate to **increase** when the cache gets larger. But in this case, with FIFO, it gets worse.




Another Simple Policy: Random

- Picks a random page to replace under memory pressure.
 - It doesn't really try to be too intelligent in picking which blocks to evict.
 - Random does depend entirely upon how lucky Random gets in its choice.

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss	3	2,0,1
2	Hit		2,0,1
1	Hit		2,0,1

Random Performance

- Sometimes, Random is as good as optimal, achieving 6 hits on the example trace.



Using History

- Lean on the past and use history.
- Two type of historical information.

Historical Information	Meaning	Algorithms
recency	The more recently a page has been accessed, the more likely it will be accessed again	LRU
frequency	If a page has been accessed many times, It should not be replaced as it clearly has some value	LFU

Using History : LRU


- Replaces the least-recently-used page.

Reference Row										
0	1	2	0	1	3	0	3	1	2	1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		1,2,0
1	Hit		2,0,1
3	Miss	2	0,1,3
0	Hit		1,3,0
3	Hit		1,0,3
1	Hit		0,3,1
2	Miss	0	3,1,2
1	Hit		3,2,1


Workload Example : The No-Locality Workload

- Each reference is to a random page within the set of accessed pages.
 - Workload accesses 100 unique pages over time.
 - Choosing the next page to refer to at random




Workload Example : The 80-20 Workload

- Exhibits locality: 80% of the **reference** are made to 20% of the pages
- Remaining 20% of the **reference** are made to the remaining 80% of the pages



Workload Example: The Looping Sequential

- Refer to 50 pages in sequence.
 - Starting at 0, then 1, ... up to page 49, and then we Loop, repeating those accesses, for total of 10,000 accesses to 50 unique pages.



Implementing Historical Algorithms


- To keep track of which pages have been least-and-recently used, the system has to do some accounting work on **every memory reference.**
 - Add a little bit of hardware support.

Approximating LRU

- Require some hardware support, in the form of a use bit
 - Whenever a page is referenced, the use bit is set by hardware to 1.
 - Hardware never clears the bit, though; that is the responsibility of the OS
- Clock Algorithm
 - All pages of the system arrange in a circular list.
 - A clock hand points to some particular page to begin with.

Clock Algorithm


- The algorithm continues until it finds a use bit that is set to 0.



When a page fault occurs, the page the hand is pointing to is inspected.
The action taken depends on the Use bit

Workload with Clock Algorithm

- Clock algorithm doesn't do as well as perfect LRU, it does better than approach that don't consider history at all.



Considering Dirty Pages



- The hardware include a modified bit (a.k.a dirty bit)
 - Page has been modified and is thus dirty, it must be written back to disk to evict it.
 - Page has not been modified, the eviction is free.

Page Selection Policy

- The OS has to decide when to bring a page into memory.
- Presents the OS with some different options.

Prefetching


- The OS guess that a page is about to be used, and thus bring it in ahead of time.



Page 2 likely soon be accessed and
thus should be brought into memory too


Clustering, Grouping

- Collect a number of pending writes together in memory and write them to disk in one write.
 - Perform a single large write more efficiently than many small ones.



Thrashing

- Memory is **oversubscribed** and the memory demands of the set of running processes **exceeds** the available physical memory.
 - Decide not to run a subset of processes.
 - Reduced set of processes working sets fit in memory.



Programming Assignment 2

Quiz 3

Homework 3

Project Launch

What's Next

- Homework Assignment 2
 - Due by Monday, February 9 by 11:59pm ET
- Quiz 2
 - Due by Monday, February 9 by 11:59pm ET
- Homework Assignment 3
 - Due by Sunday, February 15 by 11:59pm ET
- Quiz 3
 - Due by Sunday, February 15 by 11:59pm ET
- Programming Assignment 2 (PA2)
 - Due by Friday, February 27, 2026 by 11:59pm ET
- Course Project
 - Due by Friday, February 27, 2026 by 11:59pm ET